

Queries

Querying data

- **SELECT** retrieves rows from zero or more tables.
- You must have SELECT privilege on each column used in a SELECT command.

SELECT syntax

```
SELECT [ ALL | DISTINCT [ ON ( expression [, ...] ) ] ]  
      [ * | expression [ [ AS ] output_name ] [, ...] ]  
      [ FROM from_item [, ...] ]  
      [ WHERE condition ]  
      [ GROUP BY grouping_element [, ...] ]  
      [ HAVING condition [, ...] ]  
      [ { UNION | INTERSECT | EXCEPT } [ ALL | DISTINCT ] select ]  
      [ ORDER BY expression [ ASC | DESC | USING operator ] [ NULLS { FIRST | LAST } ] [, ...] ]  
      [ LIMIT { count | ALL } ]  
      [ OFFSET start [ ROW | ROWS ] ]
```

SELECT List

- The SELECT list specifies expressions that form the output rows of the SELECT statement.

SELECT List

- The SELECT list specifies expressions that form the output rows of the SELECT statement.
- The expressions can (and usually do) refer to columns computed in the FROM clause.

SELECT List

- The SELECT list specifies expressions that form the output rows of the SELECT statement.
- The expressions can (and usually do) refer to columns computed in the FROM clause.
- Just as in a table, every output column of a SELECT has a name.

SELECT List

- The SELECT list specifies expressions that form the output rows of the SELECT statement.
- The expressions can (and usually do) refer to columns computed in the FROM clause.
- Just as in a table, every output column of a SELECT has a name.
- To specify the name to use for an output column, write AS *output name* after the column's expression.

SELECT List

- If you do not specify a column name, a name is chosen automatically by PostgreSQL.

SELECT List

- If you do not specify a column name, a name is chosen automatically by PostgreSQL.
- If the column's expression is a simple column reference then the chosen name is the same as that column's name.

SELECT List

- If you do not specify a column name, a name is chosen automatically by PostgreSQL.
- If the column's expression is a simple column reference then the chosen name is the same as that column's name.
- In more complex cases a function or type name may be used, or the system may fall back on a generated name such as ?column?.

SELECT List

```
SELECT lower('HELLO'), upper('hello')
```

	lower text	upper text
1	hello	HELLO

SELECT List

SELECT 1+3, 3*4

	?column? integer	?column? integer
1	4	12

ALL vs DISTINCT

- If `SELECT DISTINCT` is specified, all duplicate rows are removed from the result set

ALL vs DISTINCT

- If `SELECT DISTINCT` is specified, all duplicate rows are removed from the result set
- One row is kept from each group of duplicates

ALL vs DISTINCT

- If `SELECT DISTINCT` is specified, all duplicate rows are removed from the result set
- One row is kept from each group of duplicates
- `SELECT ALL` specifies the opposite: all rows are kept; that is the default.

SELECT ALL * FROM cd.facilities

	facid [PK] integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	0	Tennis Court 1	5	25	10000	200
2	1	Tennis Court 2	5	25	8000	200
3	2	Badminton Court	0	15.5	4000	50
4	3	Table Tennis	0	5	320	10
5	4	Massage Room 1	35	80	4000	3000
6	5	Massage Room 2	35	80	4000	3000
7	6	Squash Court	3.5	17.5	5000	80
8	7	Snooker Table	0	5	450	15
9	8	Pool Table	0	5	400	15

SELECT DISTINCT * FROM cd.facilities

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	5	Massage Room 2	35	80	4000	3000
2	8	Pool Table	0	5	400	15
3	3	Table Tennis	0	5	320	10
4	1	Tennis Court 2	5	25	8000	200
5	4	Massage Room 1	35	80	4000	3000
6	7	Snooker Table	0	5	450	15
7	2	Badminton Court	0	15.5	4000	50
8	0	Tennis Court 1	5	25	10000	200
9	6	Squash Court	3.5	17.5	5000	80

SELECT DISTINCT ON(membercost) * FROM cd.facilities

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	7	Snooker Table	0	5	450	15
2	6	Squash Court	3.5	17.5	5000	80
3	0	Tennis Court 1	5	25	10000	200
4	5	Massage Room 2	35	80	4000	3000

WHERE clause

- The optional WHERE clause has the general form

WHERE condition

WHERE clause

- The optional WHERE clause has the general form

WHERE *condition*

- Where *condition* is any expression that evaluates to a result of type boolean.

WHERE clause

- The optional WHERE clause has the general form

WHERE *condition*

- Where *condition* is any expression that evaluates to a result of type boolean.
- Any row that does not satisfy this condition will be eliminated from the output.

```
SELECT * FROM cd.facilities WHERE membercost > 5
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	4	Massage Room 1	35	80	4000	3000
2	5	Massage Room 2	35	80	4000	3000

GROUP BY

- The optional GROUP BY clause has the general form

```
GROUP BY grouping_element [, ...]
```

GROUP BY

- The optional GROUP BY clause has the general form

`GROUP BY grouping_element [, ...]`

- GROUP BY will condense into a single row all selected rows that share the same values for the grouped expressions.

GROUP BY

- The optional GROUP BY clause has the general form

`GROUP BY grouping_element [, ...]`

- GROUP BY will condense into a single row all selected rows that share the same values for the grouped expressions.
- An *expression* used inside a *grouping_element* can be an input column name, or the name or ordinal number of an output column, or an arbitrary expression formed from input-column values.

SELECT membercost FROM cd.facilities GROUP BY membercost

	facid [PK] integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	0	Tennis Court 1	5	25	10000	200
2	1	Tennis Court 2	5	25	8000	200
3	2	Badminton Court	0	15.5	4000	50
4	3	Table Tennis	0	5	320	10
5	4	Massage Room 1	35	80	4000	3000
6	5	Massage Room 2	35	80	4000	3000
7	6	Squash Court	3.5	17.5	5000	80
8	7	Snooker Table	0	5	450	15
9	8	Pool Table	0	5	400	15

	membercost numeric
1	3.5
2	35
3	5
4	0

```
SELECT membercost, sum(guestcost) AS guest_sum
FROM cd.facilities
GROUP BY membercost
```

	facid [PK] integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	0	Tennis Court 1	5	25	10000	200
2	1	Tennis Court 2	5	25	8000	200
3	2	Badminton Court	0	15.5	4000	50
4	3	Table Tennis	0	5	320	10
5	4	Massage Room 1	35	80	4000	3000
6	5	Massage Room 2	35	80	4000	3000
7	6	Squash Court	3.5	17.5	5000	80
8	7	Snooker Table	0	5	450	15
9	8	Pool Table	0	5	400	15

	membercost numeric	guest_sum numeric
1	3.5	17.5
2	35	160
3	5	50
4	0	30.5

HAVING

- The optional HAVING clause has the general form

HAVING condition

HAVING

- The optional HAVING clause has the general form

HAVING *condition*

- where *condition* is the same as specified for the WHERE clause.

HAVING

- The optional HAVING clause has the general form

HAVING *condition*

- where *condition* is the same as specified for the WHERE clause.
- HAVING eliminates group rows that do not satisfy the condition.

HAVING

- HAVING is different from WHERE:
- WHERE filters individual rows before the application of GROUP BY

HAVING

- HAVING is different from WHERE:
- WHERE filters individual rows before the application of GROUP BY
- HAVING filters group rows created by GROUP BY

HAVING

- HAVING is different from WHERE:
- WHERE filters individual rows before the application of GROUP BY
- HAVING filters group rows created by GROUP BY
- Each column referenced in *condition* must unambiguously reference a grouping column, unless the reference appears within an aggregate function

```

SELECT membercost, sum(guestcost) AS guest_sum
FROM cd.facilities
GROUP BY membercost
HAVING membercost > 0

```

	membercost numeric	guest_sum numeric
1	3.5	17.5
2	35	160
3	5	50
4	0	30.5

	membercost numeric	guest_sum numeric
1	3.5	17.5
2	35	160
3	5	50

```

SELECT membercost, sum(guestcost)
  FROM cd.facilities
 GROUP BY membercost
HAVING sum(initialoutlay) > 5000

```

	membercost numeric	guest_sum numeric
1	3.5	17.5
2	35	160
3	5	50
4	0	30.5

	membercost numeric	guest_sum numeric
1	35	160
2	5	50
3	0	30.5

UNION

- The UNION clause has this general form:

```
select_statement UNION [ ALL | DISTINCT ] select_statement
```

UNION

- The UNION clause has this general form:

```
select_statement UNION [ ALL | DISTINCT ] select_statement
```

- *select_statement* is any SELECT statement without an ORDER BY, LIMIT clause

UNION

- The UNION clause has this general form:

```
select_statement UNION [ ALL | DISTINCT ] select_statement
```

- *select_statement* is any SELECT statement without an ORDER BY, LIMIT clause
- The UNION operator computes the set union of the rows returned by the involved SELECT statements.

UNION

- A row is in the set union of two result sets if it appears in at least one of the result sets.

UNION

- A row is in the set union of two result sets if it appears in at least one of the result sets.
- The two SELECT statements that represent the direct operands of the UNION must produce the same number of columns

UNION

- A row is in the set union of two result sets if it appears in at least one of the result sets.
- The two SELECT statements that represent the direct operands of the UNION must produce the same number of columns
- Corresponding columns must be of compatible data types.

```
SELECT * FROM cd.facilities WHERE guestcost >25
UNION SELECT * FROM cd.facilities WHERE membercost > 0
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	4	Massage Room 1	35	80	4000	3000
2	0	Tennis Court 1	5	25	10000	200
3	5	Massage Room 2	35	80	4000	3000
4	6	Squash Court	3.5	17.5	5000	80
5	1	Tennis Court 2	5	25	8000	200

```
SELECT * FROM cd.facilities WHERE guestcost >25
UNION DISTINCT SELECT * FROM cd.facilities WHERE membercost > 0
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	4	Massage Room 1	35	80	4000	3000
2	0	Tennis Court 1	5	25	10000	200
3	5	Massage Room 2	35	80	4000	3000
4	6	Squash Court	3.5	17.5	5000	80
5	1	Tennis Court 2	5	25	8000	200

```
SELECT * FROM cd.facilities WHERE guestcost > 25
UNION ALL SELECT * FROM cd.facilities WHERE membercost > 0
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	4	Massage Room 1	35	80	4000	3000
2	5	Massage Room 2	35	80	4000	3000
3	0	Tennis Court 1	5	25	10000	200
4	1	Tennis Court 2	5	25	8000	200
5	4	Massage Room 1	35	80	4000	3000
6	5	Massage Room 2	35	80	4000	3000
7	6	Squash Court	3.5	17.5	5000	80

```
SELECT * FROM cd.facilities
UNION ALL SELECT * FROM cd.facilities
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	0	Tennis Court 1	5	25	10000	200
2	1	Tennis Court 2	5	25	8000	200
3	2	Badminton Court	0	15.5	4000	50
4	3	Table Tennis	0	5	320	10
5	4	Massage Room 1	35	80	4000	3000
6	5	Massage Room 2	35	80	4000	3000
7	6	Squash Court	3.5	17.5	5000	80
8	7	Snooker Table	0	5	450	15
9	8	Pool Table	0	5	400	15
10	0	Tennis Court 1	5	25	10000	200
11	1	Tennis Court 2	5	25	8000	200
12	2	Badminton Court	0	15.5	4000	50
13	3	Table Tennis	0	5	320	10
14	4	Massage Room 1	35	80	4000	3000
15	5	Massage Room 2	35	80	4000	3000
16	6	Squash Court	3.5	17.5	5000	80
17	7	Snooker Table	0	5	450	15
18	8	Pool Table	0	5	400	15

INTERSECT

- The INTERSECT clause has this general form:

```
select_statement INTERSECT [ ALL | DISTINCT ] select_statement
```

INTERSECT

- The INTERSECT clause has this general form:

```
select_statement INTERSECT [ ALL | DISTINCT ] select_statement
```

- *select_statement* is any SELECT statement without an ORDER BY, LIMIT clause

INTERSECT

- The INTERSECT clause has this general form:

```
select_statement INTERSECT [ ALL | DISTINCT ] select_statement
```

- *select_statement* is any SELECT statement without an ORDER BY, LIMIT clause
- The INTERSECT operator computes the set intersection of the rows returned by the involved SELECT statements.

INTERSECT

- A row is in the intersection of two result sets if it appears in both result sets.

INTERSECT

- A row is in the intersection of two result sets if it appears in both result sets.
- The result of `INTERSECT` does not contain any duplicate rows unless the `ALL` option is specified.

INTERSECT

- A row is in the intersection of two result sets if it appears in both result sets.
- The result of `INTERSECT` does not contain any duplicate rows unless the `ALL` option is specified.
- With `ALL`, a row that has m duplicates in the left table and n duplicates in the right table will appear $\min(m, n)$ times in the result set.

INTERSECT

- INTERSECT binds more tightly than UNION.
- A UNION B INTERSECT C will be read as A UNION (B INTERSECT C)

```
SELECT * FROM cd.facilities WHERE guestcost > 25  
INTERSECT SELECT * FROM cd.facilities WHERE membercost > 0
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	5	Massage Room 2	35	80	4000	3000
2	4	Massage Room 1	35	80	4000	3000

```
SELECT * FROM cd.facilities WHERE guestcost >25
INTERSECT DISTINCT SELECT * FROM cd.facilities WHERE membercost > 0
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	5	Massage Room 2	35	80	4000	3000
2	4	Massage Room 1	35	80	4000	3000

```
SELECT * FROM cd.facilities WHERE guestcost >25
INTERSECT ALL SELECT * FROM cd.facilities WHERE membercost > 0
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	5	Massage Room 2	35	80	4000	3000
2	4	Massage Room 1	35	80	4000	3000

```
SELECT * FROM cd.facilities
INTERSECT SELECT * FROM cd.facilities
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	5	Massage Room 2	35	80	4000	3000
2	8	Pool Table	0	5	400	15
3	3	Table Tennis	0	5	320	10
4	1	Tennis Court 2	5	25	8000	200
5	4	Massage Room 1	35	80	4000	3000
6	7	Snooker Table	0	5	450	15
7	2	Badminton Court	0	15.5	4000	50
8	0	Tennis Court 1	5	25	10000	200
9	6	Squash Court	3.5	17.5	5000	80

EXCEPT

- The EXCEPT clause has this general form:

```
select_statement EXCEPT [ ALL | DISTINCT ] select_statement
```

EXCEPT

- The EXCEPT clause has this general form:

```
select_statement EXCEPT [ ALL | DISTINCT ] select_statement
```

- *select_statement* is any SELECT statement without an ORDER BY, LIMIT clause

EXCEPT

- The EXCEPT clause has this general form:

```
select_statement EXCEPT [ ALL | DISTINCT ] select_statement
```

- *select_statement* is any SELECT statement without an ORDER BY, LIMIT clause
- The EXCEPT operator computes the set of rows that are in the result of the left SELECT statement but not in the result of the right one.

EXCEPT

- The result of EXCEPT does not contain any duplicate rows unless the ALL option is specified.

EXCEPT

- The result of EXCEPT does not contain any duplicate rows unless the ALL option is specified.
- With ALL, a row that has m duplicates in the left table and n duplicates in the right table will appear $\max(m-n, 0)$ times in the result set.

EXCEPT

- The result of EXCEPT does not contain any duplicate rows unless the ALL option is specified.
- With ALL, a row that has m duplicates in the left table and n duplicates in the right table will appear $\max(m-n, 0)$ times in the result set.
- Multiple EXCEPT operators in the same SELECT statement are evaluated left to right, unless parentheses dictate otherwise.

EXCEPT

- The result of EXCEPT does not contain any duplicate rows unless the ALL option is specified.
- With ALL, a row that has m duplicates in the left table and n duplicates in the right table will appear $\max(m-n, 0)$ times in the result set.
- Multiple EXCEPT operators in the same SELECT statement are evaluated left to right, unless parentheses dictate otherwise.
- EXCEPT binds at the same level as UNION.

```
SELECT * FROM cd.facilities WHERE guestcost > 10
EXCEPT SELECT * FROM cd.facilities WHERE membercost > 5
```

	facid integer	name character varying (100)	membercost numeric	guestcost numeric	initialoutlay numeric	monthlymaintenance numeric
1	6	Squash Court	3.5	17.5	5000	80
2	1	Tennis Court 2	5	25	8000	200
3	2	Badminton Court	0	15.5	4000	50
4	0	Tennis Court 1	5	25	10000	200

Questions?